

Exercise-2

April 10, 2024

1 DECLAN TREVOR KINTU - 100944330

AIDI 2000 - EXERCISE 2 PROFESSOR SAJEEVA SALGADOE

1) Import libraries.

```
[ ]: import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import Xception
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
from tensorflow.keras.models import Sequential
import matplotlib.pyplot as plt
```

2) Load and preprocess data.

```
[ ]: # Load and preprocess data
train_datagen = ImageDataGenerator(rescale=1./255, validation_split=0.2)
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    r'images\train',
    target_size=(299, 299), # Xception requires input images to be at least
    ↪ 299x299
    batch_size=32,
    class_mode='binary', # or 'categorical' if more than two classes
    subset='training'
)

validation_generator = train_datagen.flow_from_directory(
    r'images\train',
    target_size=(299, 299),
    batch_size=32,
    class_mode='binary',
    subset='validation'
)
```

Found 80 images belonging to 2 classes.

Found 20 images belonging to 2 classes.

3) Define graph function.

```
[ ]: # Define function to plot performance metrics
def plot_performance(history, title):
    plt.figure(figsize=(12, 6))

    # Plot training & validation accuracy values
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('Model Accuracy ' + title)
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Validation'], loc='upper left')

    # Plot training & validation loss values
    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Model Loss ' + title)
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Validation'], loc='upper left')

    plt.tight_layout()
    plt.show()
```

4) Define and Train model.

```
[ ]: # Define the Xception model
base_model = Xception(weights='imagenet', include_top=False, input_shape=(299, 299, 3))

# Freeze the convolutional base
base_model.trainable = False

# Add classification head
model = Sequential([
    base_model,
    GlobalAveragePooling2D(),
    Dense(512, activation='relu'), # Add a fully connected layer with 512
    units and ReLU activation
    Dropout(0.5), # Add a dropout layer with dropout rate of 0.5 to reduce
    overfitting
    Dense(256, activation='relu'), # Add another fully connected layer with
    256 units and ReLU activation
    Dropout(0.5), # Add another dropout layer
    Dense(1, activation='sigmoid') # Final output layer for binary
    classification
])
```

```

])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])

# Print model summary
model.summary()

# Train the model
history = model.fit(train_generator, epochs=10,
    ↪validation_data=validation_generator)

```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
exception (Functional)	(None, 10, 10, 2048)	20861480
global_average_pooling2d_6 (GlobalAveragePooling2D)	(None, 2048)	0

Layer (type)	Output Shape	Param #
exception (Functional)	(None, 10, 10, 2048)	20861480
global_average_pooling2d_6 (GlobalAveragePooling2D)	(None, 2048)	0
dense_18 (Dense)	(None, 512)	1049088
dropout_12 (Dropout)	(None, 512)	0
dense_19 (Dense)	(None, 256)	131328
dropout_13 (Dropout)	(None, 256)	0
dense_20 (Dense)	(None, 1)	257

```

=====
Total params: 22042153 (84.08 MB)
Trainable params: 1180673 (4.50 MB)
Non-trainable params: 20861480 (79.58 MB)

```

```

-----
Epoch 1/10
3/3 [=====] - 18s 5s/step - loss: 0.6577 - accuracy:
0.5625 - val_loss: 0.4231 - val_accuracy: 0.9500

```

Epoch 2/10
 3/3 [=====] - 14s 5s/step - loss: 0.3251 - accuracy: 0.9250 - val_loss: 0.2081 - val_accuracy: 0.9500
 Epoch 3/10
 3/3 [=====] - 14s 6s/step - loss: 0.1542 - accuracy: 1.0000 - val_loss: 0.1043 - val_accuracy: 0.9500
 Epoch 4/10
 3/3 [=====] - 14s 6s/step - loss: 0.0831 - accuracy: 0.9500 - val_loss: 0.0948 - val_accuracy: 0.9500
 Epoch 5/10
 3/3 [=====] - 14s 5s/step - loss: 0.0432 - accuracy: 0.9875 - val_loss: 0.0996 - val_accuracy: 1.0000
 Epoch 6/10
 3/3 [=====] - 14s 5s/step - loss: 0.0558 - accuracy: 0.9750 - val_loss: 0.0777 - val_accuracy: 0.9500
 Epoch 7/10
 3/3 [=====] - 14s 5s/step - loss: 0.0048 - accuracy: 1.0000 - val_loss: 0.1121 - val_accuracy: 0.9500
 Epoch 8/10
 3/3 [=====] - 14s 6s/step - loss: 0.0030 - accuracy: 1.0000 - val_loss: 0.1729 - val_accuracy: 0.9500
 Epoch 9/10
 3/3 [=====] - 16s 6s/step - loss: 0.0075 - accuracy: 1.0000 - val_loss: 0.1852 - val_accuracy: 0.9500
 Epoch 10/10
 3/3 [=====] - 16s 7s/step - loss: 0.0012 - accuracy: 1.0000 - val_loss: 0.1733 - val_accuracy: 0.9500

5) Evaluate the model.

```
[ ]: # Evaluate the model
test_generator = test_datagen.flow_from_directory(
    r'images\test',
    target_size=(299, 299),
    batch_size=32,
    class_mode='binary'
)

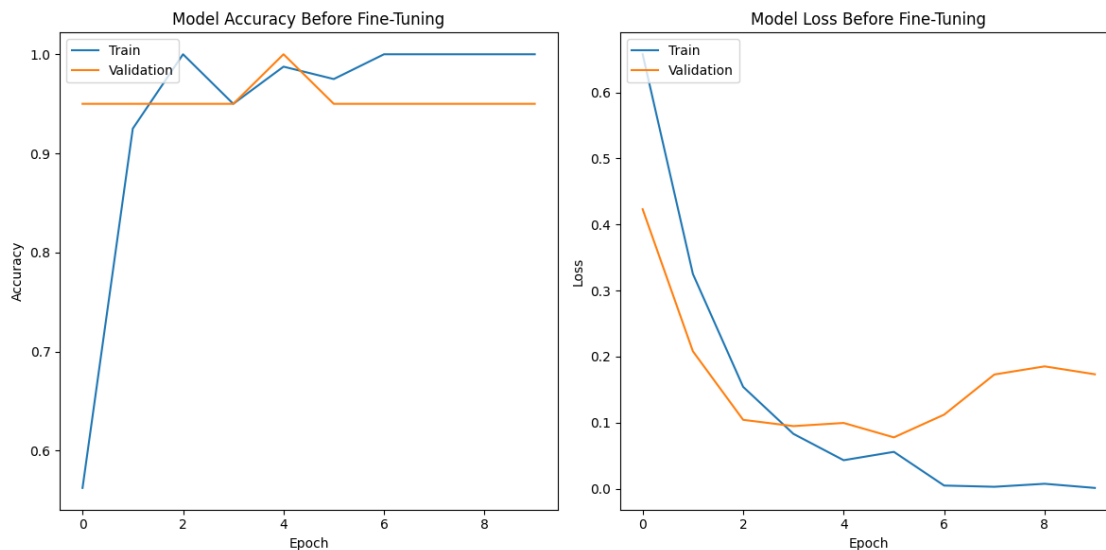
test_loss, test_acc = model.evaluate(test_generator)
print('Test accuracy:', test_acc)
```

Found 10 images belonging to 2 classes.

1/1 [=====] - 2s 2s/step - loss: 0.0617 - accuracy: 1.0000
 1/1 [=====] - 2s 2s/step - loss: 0.0617 - accuracy: 1.0000
 Test accuracy: 1.0

6) Plot model performance.

```
[ ]: plot_performance(history, 'Before Fine-Tuning')
```



7) Fine-tune the model.

```
[ ]: # Unfreeze some layers of the base model
base_model.trainable = True

# It's important to recompile the model after unfreezing the base model
model.compile(optimizer=tf.keras.optimizers.Adam(1e-5), # Lower learning rate
    ↪for fine-tuning
    loss='binary_crossentropy',
    metrics=['accuracy'])

history_fine = model.fit(train_generator, epochs=5,
    ↪validation_data=validation_generator)
```

Epoch 1/5

3/3 [=====] - 71s 18s/step - loss: 0.0432 - accuracy: 0.9875 - val_loss: 0.1676 - val_accuracy: 0.9500

Epoch 2/5

3/3 [=====] - 59s 19s/step - loss: 0.0060 - accuracy: 1.0000 - val_loss: 0.1651 - val_accuracy: 0.9500

Epoch 3/5

3/3 [=====] - 55s 17s/step - loss: 0.0098 - accuracy: 1.0000 - val_loss: 0.1627 - val_accuracy: 0.9500

Epoch 4/5

3/3 [=====] - 54s 17s/step - loss: 0.0054 - accuracy: 1.0000 - val_loss: 0.1619 - val_accuracy: 0.9500

Epoch 5/5

3/3 [=====] - 54s 17s/step - loss: 0.0607 - accuracy: 0.9875 - val_loss: 0.1600 - val_accuracy: 0.9500

8) Evaluate Fine-tuned model.

```
[ ]: # Evaluate the model
test_generator = test_datagen.flow_from_directory(
    r'images\test',
    target_size=(299, 299),
    batch_size=32,
    class_mode='binary'
)

test_loss, test_acc = model.evaluate(test_generator)
print('Test accuracy:', test_acc)
```

Found 10 images belonging to 2 classes.

1/1 [=====] - 2s 2s/step - loss: 0.0555 - accuracy: 1.0000

1/1 [=====] - 2s 2s/step - loss: 0.0555 - accuracy: 1.0000

Test accuracy: 1.0

9) Plot model performance

```
[ ]: plot_performance(history_fine, 'After Fine-Tuning')
```

